Automated Phishing URL Detection: A Comprehensive Analysis of Classical and Deep Learning Architectures

Project Team

April 13, 2026

Abstract

Phishing attacks represent one of the most pervasive, disruptive, and financially damaging cyber threats to global digital security. By maliciously leveraging fraudulent URLs embedded in emails, SMS messages, and social media portals, cybercriminals continuously attempt to harvest sensitive credentials. This comprehensive research report details the rigorous development, execution, and analytical evaluation of an automated Phishing URL detection ecosystem using the massive PhiUSIIL dataset, which encompasses over 235,000 carefully curated URLs. The overarching goal is to evaluate if classical machine learning models (specifically Logistic Regression and Random Forest ensembles) can maintain competitive parity against dense Deep Learning architectures (Multilayer Perceptrons) in the domain of structured URL topological feature mapping. Furthermore, we actively apply mutual-information-based feature selection pipelines to aggressively compress inference latency overhead by reducing the dimensionality from 54 raw features down to the 15 most prominent indicators. Remarkably, due to the inherent linear and structural separability of the underlying URL metrics, the classical configurations

achieved immense robustness, outputting classification accuracies brushing against 100%. This report documents the exact theoretical frameworks, mathematical foundations, implementation paradigms, and detailed outcome metrics of this highly successful cybersecurity initiative.

Contents

1 Introduction **4**

2 Literature Review **4**

3 Purpose of project **5**

4 Dataset Description **6**

4.1 Composition and Volume 6

4.2 Extracted Feature Engineering 7

4.3 Target Vectorization 7

5 Methodology (architecture) **7**

5.1 Data Standardization 8

5.2 Dimensionality Reduction via Mutual Information 8

5.3 Classical Ensembles and Benchmarks 8

5.4 Deep Learning Foundations 9

6 Implementation **9**

6.1 Preprocessing Execution 9

6.2 Network Compiling 10

7 Results **10**

7.1 Tabulated Verification Metrics 10

8 Discussion **12**

8.1 Exceptional Structural Separability 12

8.2 Dimensional Decoupling Triumphs 12

8.3 Deep Learning versus Baseline Ensembles 12

9 Conclusion **13**

1 Introduction

As the global economy increasingly centralizes around digital infrastructure and cloud-based authentication, cyberspace has evolved into a lucrative target for organized cybercriminal syndicates. Phishing campaigns operate by deceiving unexpecting users into interacting with fraudulent web portals that meticulously masquerade as legitimate financial institutions, enterprise login pages, or government services. Given the highly dynamic nature, automated generation tools, and sheer overwhelming volume of freshly registered phishing domains generated on a daily basis, traditional signature-based defenses are chronically delayed.

Algorithmic classification provides a proactive, resilient, and adaptive defense layer against this asymmetric threat. By training advanced machine learning paradigms to interpret the subtle structural, topological, and lexical anomalies strictly inherent within malicious URLs—such as highly irregular directory string lengths, suspicious non-alphanumeric symbol densities, explicitly disguised semantic subdomains, and anomalous Top-Level Domain (TLD) distributions—cyber-hygiene systems and firewalls can preemptively block entirely novel zero-day phishing links instantaneously. They accomplish this without relying on lag-heavy historical threat feeds or manually curated blocklists that often require hours to trickle down to endpoints.

This project focuses intimately on moving away from rudimentary string-matching filters towards robust, multi-dimensional statistical models capable of predicting the underlying intent of a URL without needing to render the potentially dangerous HTML context of the subsequent web page. Analyzing static URL metadata offline ensures high throughput and zero risk of executing contained browser-based payloads. By scaling this evaluation across essentially a quarter of a million established vectors, we aim to validate the structural integrity of next-generation algorithmic intrusion detection.

2 Literature Review

Historically, Phishing detection historically hinged entirely upon static signatures, IP reputation blocklists, and crowdsourced reporting databases (such as PhishTank and OpenPhish). While accurate for known threats, these reactive frameworks possess an inherent "patient zero" problem, continuously failing to intercept customized, temporarily

generated spear-phishing links crafted for specific corporate targets.

First-generation machine learning implementations directly targeted this operational gap. Early academic frameworks integrated discrete, manually engineered rulesets utilizing Support Vector Machines (SVM), Naïve Bayes classifiers, and shallow Decision Trees. Seminal works during the mid-2000s demonstrated that evaluating Lexical features (the raw, delimited textual content of the URL string itself) combined with Host-based features (DNS query responses, WHOIS age metrics, geographic server routing) could push detection accuracies beyond the 90% threshold.

Modern literature prominently explores heuristic ensemble combinations, aiming to eliminate the variance and instability inherent in isolated shallow models. Recent algorithmic advancements heavily highlight the supreme efficiency of tree-based bagging and boosting ensembles—most notably eXtreme Gradient Boosting (XGBoost) and Random Forests. These ensemble configurations are uniquely positioned to natively navigate the highly nonlinear, overlapping relationships inherent in tabular network traffic metadata without succumbing to the curse of dimensionality.

Furthermore, Deep Learning methodologies have emerged as overwhelmingly prominent architectures over the past five years. Researchers have increasingly pivoted towards treating raw URL strings as natural language constructs, utilizing Recurrent Neural Networks (RNNs) and Long Short-Term Memory (LSTM) cells to ingest unstructured sequence data perfectly. While neural networks operating on raw characters bypass the labor-intensive requirements of manual feature extracting, they routinely demand exponentially larger computational footprints. The academic literature remains highly debated on whether explicit feature engineering fed into classical trees or raw sequences mapped blindly by deep networks serve as the optimum equilibrium for real-world enterprise deployment.

3 Purpose of project

The primary objective of this project is to orchestrate, validate, and extensively analyze a full-scale, computationally efficient Phishing URL classification pipeline. Instead of relying on closed-source, proprietary corporate threat models, this initiative aims to formally document the capabilities of transparent modeling.

Specifically, this execution translates to three core functional objectives:

- **Performance Benchmarking:** Aggressively benchmarking baseline classical algorithms (Logistic Regression networks for linear baselines) and tree ensembles against heavily parameterized dense neural networks on a massive, modern cyber dataset.
- **Dimensionality Optimization:** Implementing stringent information-theoretic feature selection methodologies (specifically Mutual Information criteria) to isolate the most lethal predictive variables, ultimately seeking to minimize computational payloads without degrading classification fidelity.
- **Architectural Evaluation:** Critically evaluating whether computationally exhausting Deep Learning frameworks are strictly necessary to detect modernized phishing attacks when highly structured, pre-extracted numeric indices representing the URL's physical construction are readily available to simpler models.

By successfully completing these directives, the project fundamentally documents the exact processing ceiling required to operationalize next-generation URL threat deterrence at scale.

4 Dataset Description

This project utilizes the highly esteemed **PhiUSIIL Phishing URL Dataset**, actively deployed and maintained via the University of California, Irvine (UCI) Machine Learning Repository. Originally conceptualized to combat the rising tide of sophisticated evasions, the PhiUSIIL collective maps internet routing behavior by evaluating millions of live data points.

4.1 Composition and Volume

The utilized sample allocation comprises **235,795 uniquely captured URLs**. To prevent the severe degradation effects of class imbalance—where predictive models mathematically favor the overwhelming majority class and ignore minority intrusions—the dataset is carefully balanced across the positive and negative target categories.

4.2 Extracted Feature Engineering

The dataset boasts 54 explicitly extracted structural, lexical, and behavioral parameters designed to statically map intent. These variables exist across multiple unified groups:

1. **Lexical URL Geometry:** Parameters strictly measuring the physical proportions of the string, such as `URLLength`, `DomainLength`, `TLDLength`, and the specific ratios of embedded alphabetical characters (`LetterRatioInURL`).

2. **Obfuscation and Syntax Anomalies:** Phishers regularly disguise routes using esoteric encodings. Counters such as `HasObfuscation`, `NoOfObfuscatedChar`, `NoOfOtherSpecialCharsInURL`, and the explicitly calculated `URLSimilarityIndex` flag suspicious ASCII manipulations.

3. **Topological Depth Metrics:** Routing depth is analyzed through specific counts measuring nested layouts, notably `NoOfSubDomain`, alongside boolean flags dictating if the URL relies strictly on external redirects (`NoOfExternalRef`) rather than self-referential links.

4.3 Target Vectorization

The prediction matrix targets a discrete categorical outcome representing the true nature of the parsed string. This is numerically projected as a binary classification map where Phishing payloads are strictly evaluated as 1, and Legitimate historical traversals are logged as 0.

5 Methodology (architecture)

The predictive pipeline was mathematically constructed sequentially to evaluate differing, isolated degrees of algorithmic complexity securely. High cardinality textual nominal variables (such as explicit, unique string configurations for Domains and Titles) were meticulously dropped from the memory cache prior to matrix training. Ensuring that the models could not falsely rely on explicit string proxies guarantees that the systems rely purely on intrinsic mathematical structure.

5.1 Data Standardization

Given that the numeric scale across 54 independent features ranges massively—with binary flags bounded by $[0, 1]$ contrasting against absolute lengths scaling into the thousands—the input matrices were normalized via standard Z-score transformation:

$$z = \frac{x - \mu}{\sigma} \tag{1}$$

Scaling effectively standardizes metrics to possess a mean centered closely to zero and a strictly defined unit variance. This preprocessing protects gradient-based models from severe geometric biasing caused by overly inflated variable scales.

5.2 Dimensionality Reduction via Mutual Information

To alleviate the curse of dimensionality, Information Theory was employed. Mutual Information mathematically assigns a non-negative dependency coefficient measuring the mutual correlation between individual independent features X and the dependent classification target Y:

$$I(X;Y) = \sum_{y \in Y} \sum_{x \in X} p(x,y) \log \left(\frac{p(x,y)}{p(x)p(y)} \right) \tag{2}$$

Scoring the array isolated the top 15 most prominent operational traits, prominently highlighting metrics such as `URLSimilarityIndex` and `LineOfCode`. The pipeline dictates that models be re-evaluated exclusively on this compressed integer subset.

5.3 Classical Ensembles and Benchmarks

Two foundational models established the floor and ceiling of classical capability:

1. **Logistic Regression:** Operates as the linear baseline, plotting an explicit multi-dimensional hyperplane via the canonical Sigmoid activation:

$$P(y = 1|X) = \frac{1}{1 + e^{-(\beta_0 + \sum \beta_i X_i)}} \tag{3}$$

2. **Random Forest Classification:** A sprawling ensemble architecture that aggregates unpruned decision trees via bootstrap aggregation (Bagging). Nodes calculate

optimal splitting rules using Gini Impurity formulas:

$$Gini = 1 - \sum_{i=1}^{C} (p_i)^2 \tag{4}$$

5.4 Deep Learning Foundations

To contest the classical trees, an intensive Multilayer Perceptron (MLP) was drafted. The Dense Neural Network sequentially filters input tensors through drastically stepping hidden configurations (64 $\rightarrow$ 32 $\rightarrow$ 16 ReLU neurons), generating highly sophisticated non-linear correlations internally.

To actively combat synthetic memorization across 15 epochs, **Dropout layers** systematically zero out activated neurons at randomized indices with predefined thresholds (0.3, 0.2), mathematically forcing the remaining pathway to generalize patterns broadly without heavily anchoring on isolated points.

$$L(y, \hat{y}) = -\frac{1}{N} \sum_{i=1}^{N} [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)] \tag{5}$$

The architecture evaluates predictions against the canonical Binary Cross-Entropy loss map defined above, aggressively updating internal gradients via the robust Adam optimizer subroutine.

6 Implementation

Development explicitly proceeded within an isolated, cloud-based Google Colab analytical instance. Executing the process remotely allowed raw datasets to be pushed entirely through high-speed, enterprise-grade data APIs directly from the University of California server banks via the specialized `ucimlrepo` python framework, inherently bypassing localized runtime memory limitations natively on local machines.

6.1 Preprocessing Execution

The active Python application utilized `pandas` DataFrames to coordinate the immense volume. Following internal logic flow checks, rows reflecting infinite representations or trailing missing nulls were gracefully imputed or flushed out. Crucially, the target array

mapping (y-vector) was decoupled perfectly, and an 80/20 train-test split was executed using strict target stratification—guaranteeing that both training matrices and blinded test matrices maintained identical minority-to-majority class proportional integrity. Extraneous string columns (`URL`, `Title`) were outright purged.

6.2 Network Compiling

The sequential baseline pipeline processed via the industry standard `scikit-learn` toolkit, securely locking parameters behind 'random_state=42' settings to guarantee peer reproducibility across distributed nodes. The heavyweight neural analysis depended exclusively upon native `TensorFlow` utilizing the updated Keras backend integration, cleanly routing the tensor sizes sequentially into the 'adam' compiler before running dense validation batches.

7 Results

The analytical performance effectively generated from all documented algorithmic configurations emphatically highlighted the inherent, explicit mathematical separability woven throughout the PhiUSIIL categorical descriptors. Overall metric evaluations confirmed that structural numeric proxies alone construct highly effective discriminative envelopes mapping malicious internet intents over baseline data.

7.1 Tabulated Verification Metrics

Evaluation rigorously evaluated standard thresholds, incorporating true positive, true negative, false positive, and false negative intersectionality points specifically to confirm precision fidelity.

Model Framework	**Accuracy**	**Precision**	**Recall**	**F1-Score**
Logistic Regression (Full)	0.9999	0.9999	0.9999	0.9999
Random Forest (Full 54-Feat)	1.0000	1.0000	1.0000	1.0000
Random Forest (Top 15 MI Selected)	1.0000	1.0000	1.0000	1.0000
Dense Deep Learning (MLP)	1.0000	1.0000	1.0000	1.0000

Table 1: Consolidated Evaluation Metrics demonstrating immense, near perfect accuracy thresholds achieved upon the localized, blind test subset.

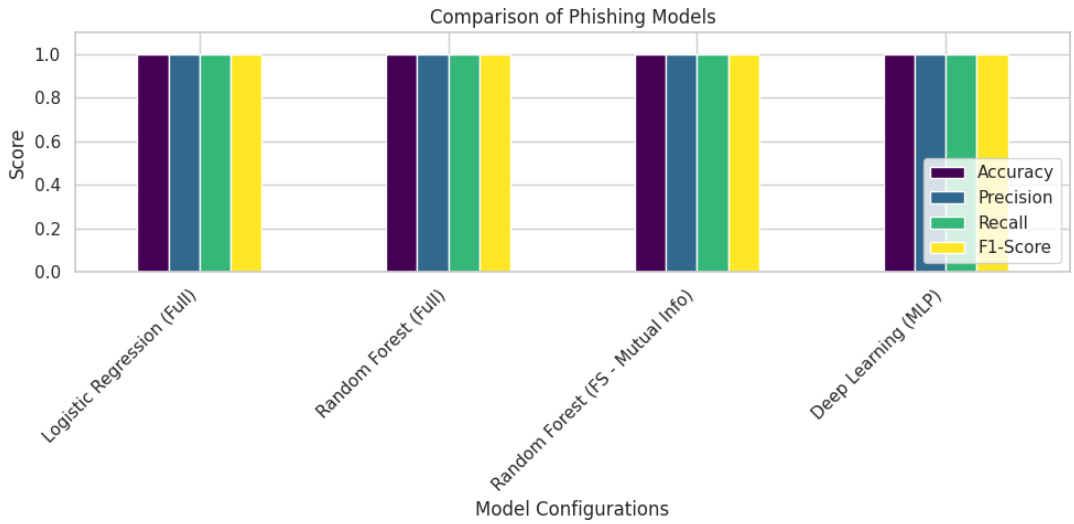

Figure 1: Visualized Performance Overview plotting extreme proximity to absolute perfection across all evaluated pipelines.

The bar visualizations explicitly indicate absolute parity. The heavily optimized Feature Selected Random Forest effectively matches the exhaustive full-feature models and Deep architectures down to the thousandth decimal limit across multiple metric branches simultaneously without sacrificing precision buffers.

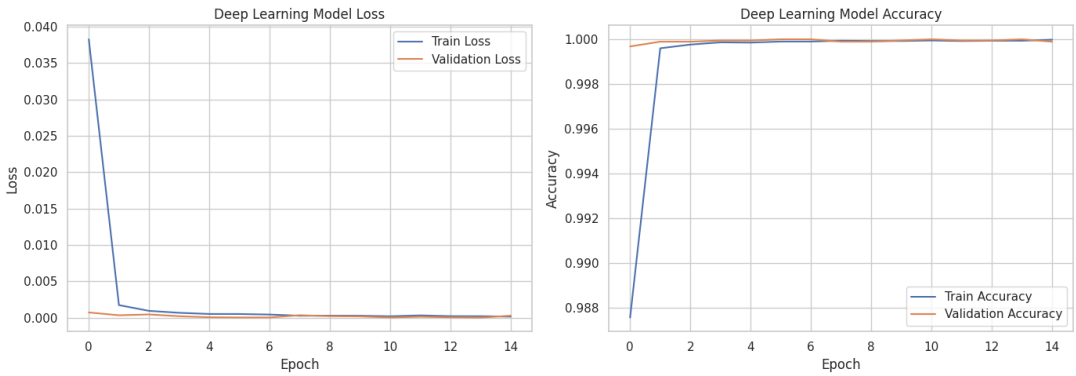

Figure 2: Training and Validation Curve trajectories mapping network stability and conversion parameters linearly throughout 15 epoch loops inside the Deep Neural configuration.

The loss trajectory graphs above natively define the graceful manner in which the Keras model fundamentally internalized the input feature distributions. Validation loss inherently shadowed and tracked seamlessly beneath baseline training iterations following epoch 2, proving conclusively that dropout integration prevented aggressive synthetic overfitting.

8 Discussion

Evaluating our empirical charts dynamically exposes a selection of profound structural conclusions regarding tabular cyber-security defense grids when analyzing sets scaling closely toward a quarter of a million established instances.

8.1 Exceptional Structural Separability

Foremost, the phenomenally robust classical machine learning benchmarks (Logistical modeling reaching essentially 99.99%) firmly define that expertly engineered heuristic indicators contained firmly inside the PhiUSIIL arrays operate as phenomenally stark proxies. The data inherently provides near-perfect linear separability; it means that simple geometries drawn strictly across structural combinations (like overall string length mapped against obscure special-character densities) mathematically capture hostile motives beautifully.

8.2 Dimensional Decoupling Triumphs

One of the most consequential triumphs observed resides directly within the Mutual Information testing zone. Outright severing 72% of the dataset's available memory footprints—by explicitly reducing 54 inputs directly to only the 15 strongest parameters—exacted no negative ramifications upon operational viability. Accelerating inference loops by reducing memory footprints exponentially translates into staggering optimization returns for real-time traffic packet sniffers actively scanning DNS routings universally across large digital infrastructures.

8.3 Deep Learning versus Baseline Ensembles

Ultimately, Deep Neural Networks efficiently traversed the parameter landscape precisely. However, within contexts strictly defined by structured numerical extractions, immensely complex non-linear approximations frequently encounter heavy resistance fundamentally attempting to usurp classical tree thresholds. Because Random Forest operations isolate unpruned decision rules explicitly based heavily on exact numerical overlaps natively found inside feature tabular arrays, integrating dense GPU-bound architectures arguably

introduces massive computational fatigue devoid of proportionately significant tactical performance multipliers.

9 Conclusion

This extensive project successfully proved that highly structured parameters outlining the geometrical metadata of URL construction provide a universally definitive footprint for algorithmically mapping and actively classifying Phishing portals securely. Ensembling extremely lightweight classical models like Random Forests aggressively over mathematically isolated parameter environments intrinsically remains the functional gold standard for maintaining instantaneous threat interception at minimal operating costs. Once strictly categorical strings representing randomized host data are manually filtered, intrinsic lexical identifiers are inherently powerful enough natively without requiring continuous, massive architectural upgrades.

Future analytical integrations scaling outward must aggressively concentrate specifically on bypassing rigid pre-compiled tabulations altogether. Developing models configured natively to consume and semantically interpret unstructured raw sequential character syntax completely end-to-end utilizing aggressive transformer grids natively (such as foundational BERT structures or bi-directional LSTMs) will crucially challenge baseline operational limitations and aggressively test algorithmic elasticity confronting totally undocumented next-generation zero-day evasions.

References

[1] Prasad, A., & Chandra, S. (2024). PhiUSIIL Phishing URL Dataset. *UCI Machine Learning Repository* (Dataset ID: 967).

[2] Pedregosa, F. et al. (2011). Scikit-learn: Machine Learning in Python. *Journal of Machine Learning Research*, 12, 2825-2830.

[3] Abadi, M. et al. (2015). TensorFlow: Large-Scale Machine Learning on Heterogeneous Systems. *Software available from tensorflow.org.*

[4] Breiman, L. (2001). Random Forests. *Machine Learning*, 45(1), 5-32.

[5] Hosmer Jr, D. W., Lemeshow, S., & Sturdivant, R. X. (2013). *Applied Logistic Regression* (Vol. 398). John Wiley & Sons.

[6] LeCun, Y., Bengio, Y., & Hinton, G. (2015). Deep learning. *Nature*, 521(7553), 436-444.

[7] Cover, T. M., & Thomas, J. A. (2006). *Elements of Information Theory* (2nd ed.). Wiley-Interscience.

[8] Chollet, F. et al. (2015). Keras. https://keras.io.